

# Linear Classification

## Machine Learning


Daniele Loiacono



**POLITECNICO**  
MILANO 1863

# References

- *Pattern Recognition and Machine Learning*, Bishop
  - ▶ Chapter 4 (4.1.1 - 4.1.3, 4.1.7, 4.3.1, 4.3.2)



# Outline

- ❑ Linear Classification Models
- ❑ Direct Approaches
  - ▶ Least Squares for Classification
  - ▶ The Perceptron Algorithm
- ❑ Probabilistic Discriminative Approach
  - ▶ Logistic Regression

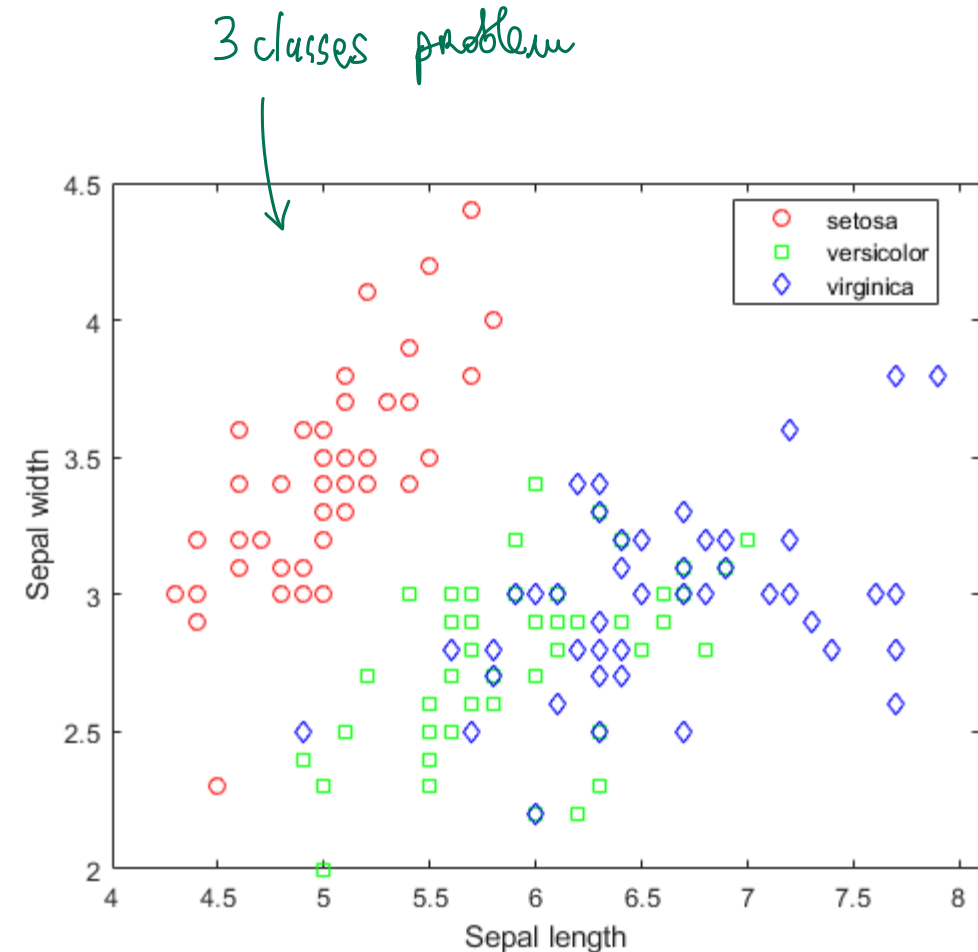
**A quick recap**

# What is classification?

- Learn, from a dataset  $\mathcal{D}$ , an **approximation** of function  $f(x)$  that maps input  $x$  to a discrete class  $C_k$  (with  $k = 1, \dots, K$ )

$$\mathcal{D} = \{\langle x, C_k \rangle\} \Rightarrow C_k = f(x)$$

- ▶ How do we model  $f$ ?
- ▶ How do we encode  $C_k$ ? ) *Model*
- ▶ How do we evaluate our approximation?
- ▶ How do we optimize our approximation?



# Classification approaches

- ❑ Discriminant function
  - ▶ Model a parametric function that maps input to classes
  - ▶ Learn parameters from data
- ❑ Probabilistic discriminative approach
  - ▶ Design a parametric model of  $p(C_k|\mathbf{x})$
  - ▶ Learn model parameters from data
- ❑ Probabilistic generative approach
  - ▶ Model  $p(\mathbf{x}|C_k)$  and class priors  $p(C_k)$
  - ▶ Fit models to data
  - ▶ Infer **posterior** with Bayes' rule:  $p(C_k|\mathbf{x}) = p(\mathbf{x}|C_k)p(C_k)/p(\mathbf{x})$

# Classification approaches

## ❑ Discriminant function

- ▶ Model a parametric function that maps input to classes
- ▶ Learn parameters from data

## ❑ Probabilistic discriminative approach

- ▶ Design a parametric model of  $p(C_k|\mathbf{x})$
- ▶ Learn model parameters from data

## ❑ Probabilistic generative approach

- ▶ Model  $p(\mathbf{x}|C_k)$  and class priors  $p(C_k)$
- ▶ Fit models to data
- ▶ Infer **posterior** with Bayes' rule:  $p(C_k|\mathbf{x}) = p(\mathbf{x}|C_k)p(C_k)/p(\mathbf{x})$

## **Discriminant Function**



# Generalized Linear Models for classification

- In linear classification, we will use <sup>GLM</sup> **generalized linear models**:

$$f(\mathbf{x}, \mathbf{w}) = f \left( w_0 + \sum_{j=1}^{D-1} w_j x_j \right) = f \left( \underbrace{\mathbf{x}^T \mathbf{w} + w_0}_{\text{our model in linear regression.}} \right) \quad f(\text{linear model})$$



- ▶  $f(\cdot)$  is **not** linear in  $\mathbf{w}$  due to the (non linear) **activation function**  $f$ , because its output is either a discrete label or a probability value
- ▶  $f(\cdot)$  partitions the input space into **decision regions** whose boundaries are called **decision boundaries or decision surfaces**
- ▶ these decision surfaces are linear function of  $\mathbf{x}$  and  $\mathbf{w}$ , as they correspond to  $\mathbf{x}^T \mathbf{w} + w_0 = \text{const}$
- ▶ generalized linear models are more complex to use with respect to linear models (both from a computational and analytical perspective)

That's why  
"GLM"

# Label Encoding

- K=2 classes*
- ❑ A common encoding for two-class problems is a binary encoding:  $t \in \{0,1\}$ 
    - ▶  $t = 1$  encodes **positive** class and  $t = 0$  encodes **negative** one
    - ▶ with this encoding,  $t$  and  $f(\cdot)$  represent the **probability** of positive class
  - ❑ A possible alternative encoding for two-class problems is  $t \in \{-1,1\}$ 
    - ▶ this encoding is convenient for some algorithms
- "Binary problem"*

- K classes*
- ❑ When the problem has  $K$  classes, a typical choice is **1-of-K** encoding
    - ▶  $t$  is a vector of length  $K$ , with a 1 in the position corresponding to the encoded class
    - ▶ with this encoding,  $t$  and  $f(\cdot)$  represent the probability density over the classes
    - ▶ as an example, a data sample that belongs to class 4 of a problem with  $K=5$ , is encoded as  $t = (0,0,0,1,0)^T$

*"Multiclass problem"*

*OR "ONE HOT ENCODING"*

# Discriminant linear function for a two-class problem

Binary classification

we use a slightly different notation compared to lin. reg.:  $w_0$  is not included in  $w$ .

Reason: for some properties we need to see explicitly  $w_0$ .

$$f(\mathbf{x}, \mathbf{w}) = \begin{cases} C_1, & \text{if } \mathbf{x}^T \mathbf{w} + w_0 \geq 0 \\ C_2, & \text{otherwise} \end{cases}$$

encoded either via  $\{0,1\}$  or via  $\{-1,1\}$ .

## Properties

► DS is  $y(\cdot) = \mathbf{x}^T \mathbf{w} + w_0 = 0$


► DS is orthogonal to  $\mathbf{w}$ :  $(\vec{x}_A - \vec{x}_B) \cdot \mathbf{w} = \vec{x}_A^T \mathbf{w} - \vec{x}_B^T \mathbf{w} = -w_0 - (-w_0) = 0$

► distance of DS from origin is  $-\frac{w_0}{\|\mathbf{w}\|_2} \left( = \frac{\vec{x}_A^T \mathbf{w}}{\|\mathbf{w}\|_2} \right)$

► "distance" of  $\mathbf{x}$  from DS is  $\frac{y(\mathbf{x})}{\|\mathbf{w}\|_2}$

"." because it has a sign.



MB: We can use that as a measure of "confidence" of our classification. The more we are far from the DS, the more confident we are in our classification.  $\rightarrow$  of later is SVM.



Now, our question is: HOW TO LEARN  $\mathbf{w}$ ?

# How to deal with multiple classes problems?

- 2 main methods
- In a multi-class problem we have  $K$  classes
  - **One-versus-the-rest** approach uses  $K-1$  binary classifiers (i.e., that solve a two-class problem)
    - ▶ each classifier discriminates  $C_i$  and not  $C_i$  regions
    - ▶ **ambiguity**: region mapped to several classes
  - **One-versus-one** approach uses  $K(K-1)/2$  class binary classifiers
    - ▶ each classifier discriminates between  $C_i$  and  $C_j$
    - ▶ similar ambiguity of previous approach



How To Solve This Issue Of Ambiguity?

# A simple solution for multiple classes

- A possible solution is to use  $K$  linear discriminant functions:

$$y_k(\mathbf{x}) = \mathbf{x}^T \mathbf{w}_k + w_{k0}, \text{ where } k = 1, \dots, K$$

- ▶ Map  $\mathbf{x}$  to class  $C_k$  if  $y_k > y_j \forall j \neq k$
- ▶ No ambiguity !!
- ▶ DS are singly connected and convex


Let  $\mathbf{x}_A, \mathbf{x}_B \in \mathcal{R}_k$

Thus,  $y_k(\mathbf{x}_A) > y_j(\mathbf{x}_A)$  and  $y_k(\mathbf{x}_B) > y_j(\mathbf{x}_B)$

$$\Rightarrow \forall \alpha (0 < \alpha < 1)$$

$$y_k(\alpha \mathbf{x}_A + (1 - \alpha) \mathbf{x}_B) > y_j(\alpha \mathbf{x}_A + (1 - \alpha) \mathbf{x}_B)$$

I choose the class in which the model is the most confident.



HOW TO FIND  $\mathbf{w}$ ?

for  
method

## Least Squares for Classification

It's possible but it has some problems.

1-hot encoding

- Let consider a K-class problem and use a 1-of-K encoding for target
- Each class is modeled with a linear function:

$$y_k(\mathbf{x}) = \mathbf{x}^T \mathbf{w}_k + w_{k0}, \text{ where } k = 1, \dots, K$$

- In matrix notation:

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$$

- ▶  $\tilde{\mathbf{W}}$  has size  $(D+1) \times K$
- ▶ k-th column of  $\tilde{\mathbf{W}}$  is  $\tilde{\mathbf{w}}_k = (w_{k0}, \mathbf{w}_k^T)^T$
- ▶  $\tilde{\mathbf{x}} = (1, \mathbf{x}^T)^T$

## Least Squares for Classification (2) *It's possible but it has some problems.*

- Given a dataset  $\mathcal{D} = \{\mathbf{x}_i, \mathbf{t}_i\}$ , where  $i=1, \dots, N$
- We can apply Least Squares to find the optimal value of  $\tilde{\mathbf{W}}$


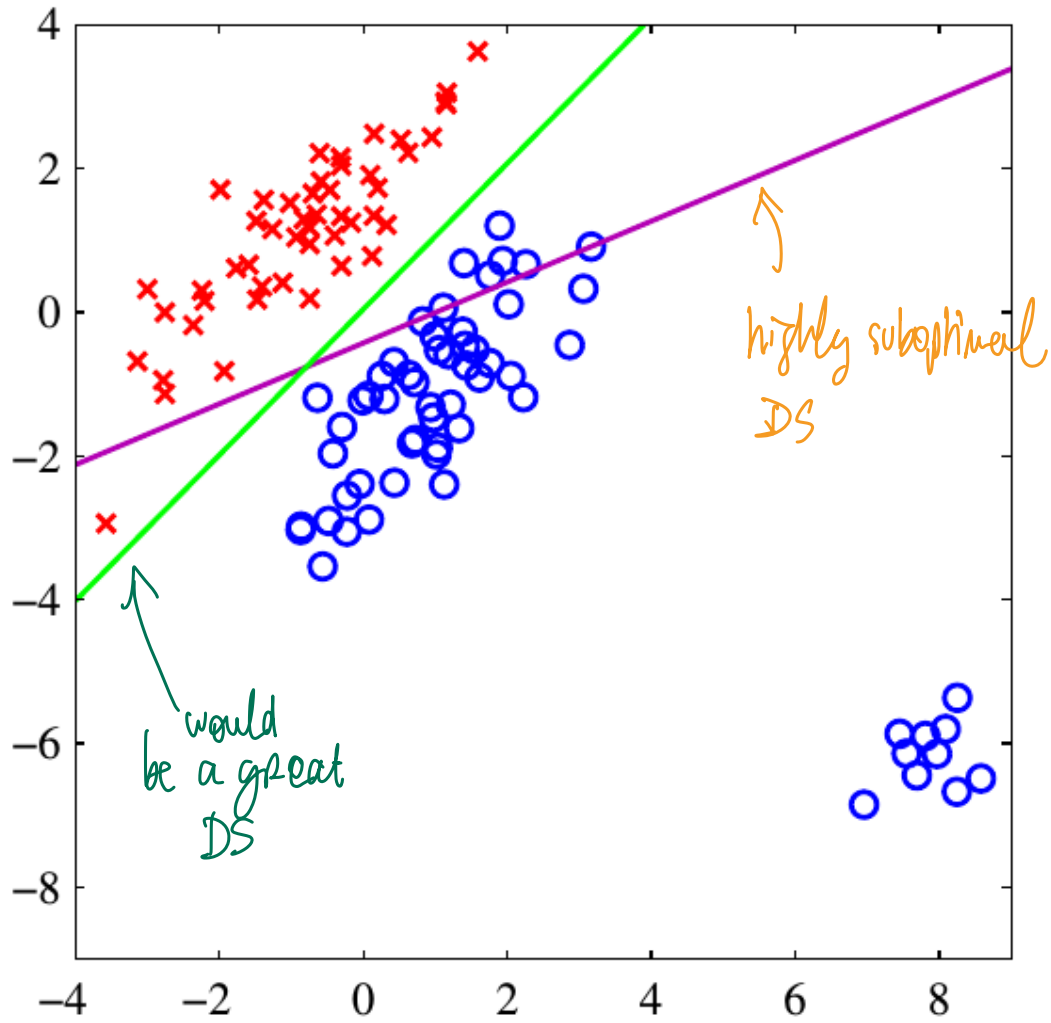
$$\tilde{\mathbf{W}} = \left( \tilde{\mathbf{X}}^T \tilde{\mathbf{X}} \right)^{-1} \tilde{\mathbf{X}}^T \mathbf{T}$$

- ▶  $\tilde{\mathbf{X}}$  is a  $N \times (D+1)$  matrix whose  $i$ -th is  $\tilde{\mathbf{x}}_i^T$
  - ▶  $\mathbf{T}$  is a  $N \times K$  matrix whose  $i$ -th row is  $\mathbf{t}_i^T$
- Any new sample  $\tilde{\mathbf{x}}_{new}^T$  is mapped to class  $C_k$  if  $t_k > t_j \quad \forall j$ , where  $t_k$  is the  $k$ -th component of the model output, computed as  $t_k = \tilde{\mathbf{x}}^T \tilde{\mathbf{w}}_k$

1st method

# Problems with Least Squares: Outliers and Output Distribution

→ OLS is **NOT** a good approach for classification problem!



WHAT IS THEN A GOOD SOLUTION?



# Linear Basis Function Models

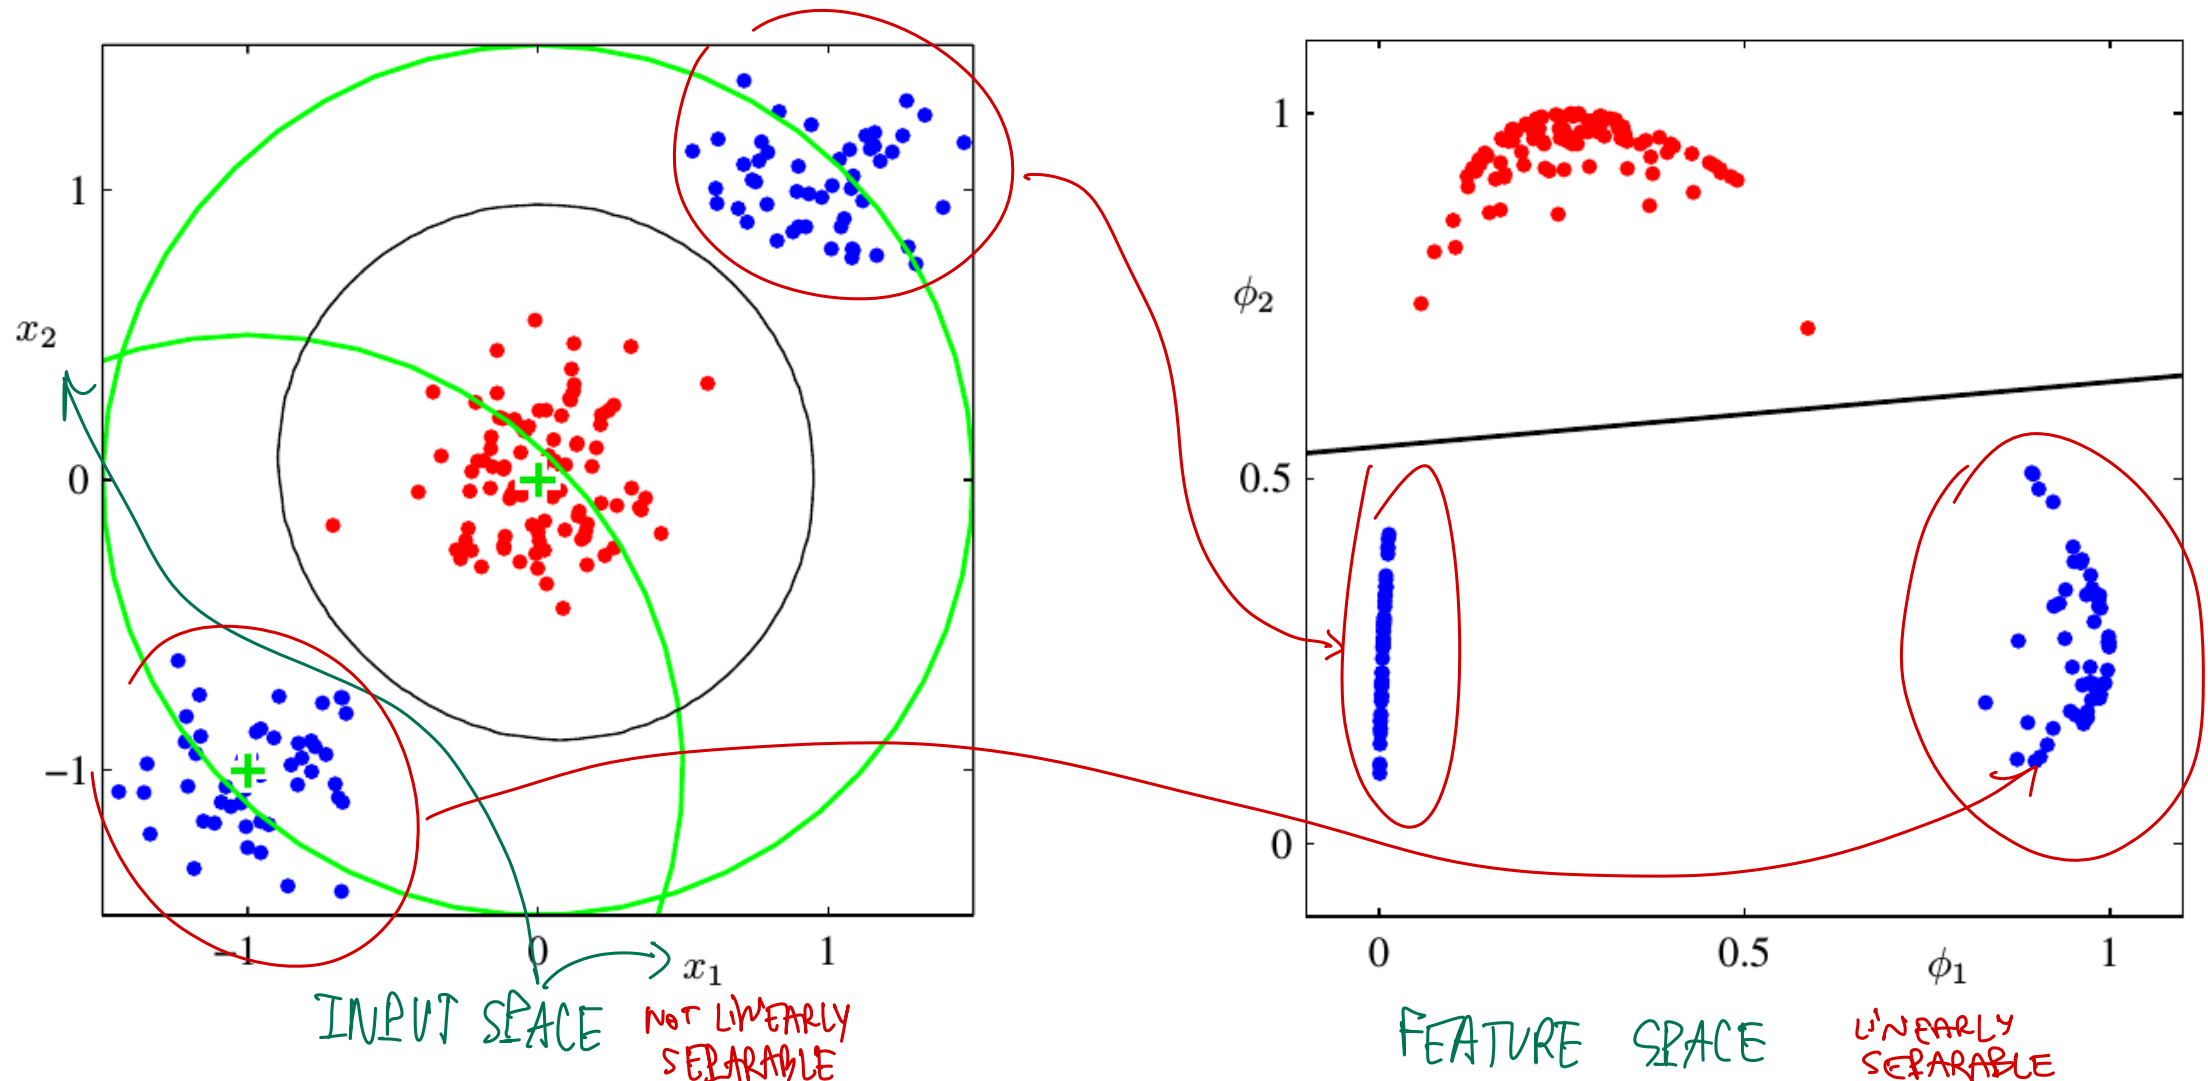
- ❑ So far we considered models that work in the **input space** (i.e., with  $\mathbf{x}$ )
- ❑ However, we can still extend models by using a fixed set of basis function  $\phi(\mathbf{x})$
- ☒ Basically, we apply a non-linear transformation to map the **input space** into a **feature space**
- ☒ As a result, decision boundaries that are **linear** in the **feature space** would correspond to **nonlinear** boundaries in the **input space**
- ☒ This allows to apply linear classification models also to problem where samples are not **linearly separable**



Paraphrase  
so3

# Linear Basis Function Models: an example

- Assuming two Gaussian basis functions (in green)



2nd  
method

## Perceptron

SO: HOW CAN WE FIND THE VECTOR  $w$ ?

- ❑ The **perceptron** is a linear discriminant model proposed by Rosenblatt in 1958 along with a **sequential learning algorithm**
- ❑ Perceptron is devised for a two-class problem, where classes encoding is  $\{-1, 1\}$

$$f(\mathbf{x}, \mathbf{w}) = \begin{cases} +1, & \text{if } \mathbf{w}^T \phi(\mathbf{x}) \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

*we use a feature space of any type.*

- ❑ The perceptron algorithm aims at finding a decision surface (also called separating hyperplane) by minimizing the distance of misclassified samples to the boundary
- ❑ Minimization of this loss function can be performed using stochastic gradient descent
- ❑ Despite a simpler loss function could be used in principle (e.g., number of misclassified samples), this are more complex to minimize.

# Perceptron Criterion

- We aim at finding  $\mathbf{w}$  such that  $\mathbf{w}^T \phi(\mathbf{x}_i) \geq 0$  for  $\mathbf{x}_i \in C_1$  and  $\mathbf{w}^T \phi(\mathbf{x}_i) < 0$  otherwise
- The Perceptron Criterion is defined as:

$$L_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \overbrace{\mathbf{w}^T \phi(\mathbf{x}_n)}^{y(\mathbf{x}_n)} t_n \rightarrow \begin{cases} \text{if misclassified} & \begin{cases} y(\mathbf{x}_n) < 0 \text{ and } t_n > 0, \\ y(\mathbf{x}_n) > 0 \text{ and } t_n < 0. \end{cases} \end{cases}$$

to make  $L_P$  positive



- ▶ samples classified correctly do not contribute to  $L$
- ▶ each misclassified sample  $\mathbf{x}_i \in \mathcal{M}$  contributes as  $\mathbf{w}^T \phi(\mathbf{x}_i) t_i$

- $L_P$  can be minimized using **stochastic gradient descent**:

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \alpha \nabla L_P(\mathbf{w}) = \mathbf{w}^{(k)} + \alpha \phi(\mathbf{x}_n) t_n$$

learning rate

- ▶ Since the scale of  $\mathbf{w}$  does not change the perceptron function, usually the **learning rate**  $\alpha$  is set to 1

2nd  
method

# Perceptron Algorithm

Given  $\mathcal{D} = \{\mathbf{x}_i, t_i\}$ , where  $i=1, \dots, N$

Initialize  $\mathbf{w}_0$

$k \leftarrow 0$

repeat

$k \leftarrow k+1$

$n \leftarrow k \bmod N$   $\nwarrow \hat{t}_n = f(\mathbf{x}_n)$

if  $\hat{t}_n \neq t_n$  then

$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \phi(\mathbf{x}_n)t_n$

endif

until convergence

otherwise: no convergence //

The algorithm works **only** if it is possible to perfectly separate the points (with no error) via a linear boundary.

2nd  
method

# Perceptron Algorithm

Given  $\mathcal{D} = \{\mathbf{x}_i, \mathbf{t}_i\}$ , where  $i=1, \dots, N$

Initialize  $\mathbf{w}_0$

$k \leftarrow 0$

repeat

$k \leftarrow k+1$


$n \leftarrow k \bmod N$

if  $\hat{t}_n \neq t_n$  then

$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \phi(\mathbf{x}_n)t_n$

endif

until convergence



2nd  
method

## Perceptron Algorithm

Given  $\mathcal{D} = \{\mathbf{x}_i, \mathbf{t}_i\}$ , where  $i=1, \dots, N$

Initialize  $\mathbf{w}_0$

$k \leftarrow 0$

**repeat**

$k \leftarrow k+1$


$n \leftarrow k \bmod N$

**if**  $\hat{t}_n \neq t_n$  **then**

$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \phi(\mathbf{x}_n)t_n$

**endif**

**until** convergence



2nd  
method

## Perceptron Algorithm

Given  $\mathcal{D} = \{\mathbf{x}_i, \mathbf{t}_i\}$ , where  $i=1, \dots, N$

Initialize  $\mathbf{w}_0$

$k \leftarrow 0$

repeat

$k \leftarrow k+1$


$n \leftarrow k \bmod N$

    if  $\hat{t}_n \neq t_n$  then

$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \phi(\mathbf{x}_n)t_n$

    endif

until convergence





2nd  
method

## Perceptron Algorithm

Given  $\mathcal{D} = \{\mathbf{x}_i, \mathbf{t}_i\}$ , where  $i=1, \dots, N$

Initialize  $\mathbf{w}_0$

$k \leftarrow 0$

**repeat**

$k \leftarrow k+1$


$n \leftarrow k \bmod N$

**if**  $\hat{t}_n \neq t_n$  **then**

$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \phi(\mathbf{x}_n)t_n$

**endif**

**until** convergence



# Perceptron Convergence Theorem

- 
- A single update **reduce the error** on the **single misclassified sample**:

$$-\mathbf{w}^{(k+1)T} \phi(\mathbf{x}_n) t_n = -\mathbf{w}^{(k)T} \phi(\mathbf{x}_n) t_n - (\phi(\mathbf{x}_n) t_n)^T \phi(\mathbf{x}_n) t_n < -\mathbf{w}^{(k)T} \phi(\mathbf{x}_n) t_n$$


- ▶ This does **not** imply that the entire loss is reduced after each update

- Perceptron Convergence Theorem:

We have no  
guaranty that  
we are pro-  
gressing @  
each step...

*If the training data set is **linearly separable** in the feature space  $\Phi$ , then the perceptron learning algorithm is guaranteed to find an **exact solution** in a **finite number of steps**.*

- ▶ How many steps? **Several steps might be necessary**, thus it might be difficult to distinguish between **nonseparable** problems and **slowly converging** ones

- 
- ▶ Which solution? **If multiple solutions exist, the one found by the algorithms depends from initialization of parameters and the order of updates**


## Probabilistic Discriminative Approaches

# Two-Class Logistic Regression

- In a discriminative approach we model directly the conditioned class probability:

$$p(C_1|\phi) = \frac{1}{1 + \exp(-\mathbf{w}^T \phi)} = \sigma(\mathbf{w}^T \phi)$$

- ▶ where  $\sigma(a) = 1/(1 + \exp(-a))$  is the sigmoidal function
- ▶  $p(C_2|\phi) = 1 - p(C_1|\phi)$
- ▶ this model is known as **logistic regression**



## Maximum Likelihood for Logistic Regression

- Given dataset  $\mathcal{D} = \{\mathbf{x}_i, t_i\}$ , where  $i=1, \dots, N$  and  $t_i \in \{0, 1\}$  we want to maximize the likelihood, i.e., the probability to observe the targets given the inputs:  $p(\mathbf{t}|\mathbf{X}, \mathbf{w})$

## Maximum Likelihood for Logistic Regression

- Given dataset  $\mathcal{D} = \{\mathbf{x}_i, t_i\}$ , where  $i=1, \dots, N$  and  $t_i \in \{0, 1\}$  we want to maximize the likelihood, i.e., the probability to observe the targets given the inputs:  $p(\mathbf{t}|\mathbf{X}, \mathbf{w})$
- We model the likelihood of the single sample using a Bernoulli distribution, using the logistic regression model for conditioned class probability:

$$p(t_n | \mathbf{x}_n, \mathbf{w}) = \underbrace{y_n^{t_n} (1 - y_n)^{1-t_n}}_{= p(t_n=1 | \mathbf{x}_n, \mathbf{w}) \text{ if } t_n = 1 \text{ \& otherwise } 1 - p(t_n=1 | \mathbf{x}_n, \mathbf{w})} \quad \text{where} \quad y_n = p(t_n = 1 | \mathbf{x}_n, \mathbf{w}) = \sigma(\mathbf{w}^T \phi_n)$$

- Assuming that data in  $\mathcal{D}$  have been independently sampled we get:

$$p(\mathbf{t} | \mathbf{X}, \mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}, \quad y_n = \sigma(\mathbf{w}^T \phi_n)$$

## Maximum Likelihood for Logistic Regression (2)

- A convenient loss function to minimize is the negative log-likelihood (also known as “cross-entropy error” function)

$$L(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{X}, \mathbf{w}) = -\sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n)) = \sum_{n=1}^N L_n$$

- Now we have to compute the derivative of  $L(\mathbf{w})$ :

$$\frac{\partial L_n}{\partial y_n} = \frac{y_n - t_n}{y_n(1 - y_n)}, \quad \frac{\partial y_n}{\partial \mathbf{w}} = y_n(1 - y_n)\phi_n,$$


## Maximum Likelihood for Logistic Regression (2)

- A convenient loss function to **minimize** is the negative log-likelihood (also known as **cross-entropy error function**)

$$L(\mathbf{w}) = -\ln p(\mathbf{t}|\mathbf{X}, \mathbf{w}) = -\sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n)) = \sum_{n=1}^N L_n$$

- Now we have to compute the derivative of  $L(\mathbf{w})$ :

$$\frac{\partial L_n}{\partial y_n} = \frac{y_n - t_n}{y_n(1 - y_n)}, \quad \frac{\partial y_n}{\partial \mathbf{w}} = y_n(1 - y_n)\phi_n, \implies \frac{\partial L_n}{\partial \mathbf{w}} = \frac{\partial L_n}{\partial y_n} \frac{\partial y_n}{\partial \mathbf{w}} = (y_n - t_n)\phi_n$$

 Chain rule



## Maximum Likelihood for Logistic Regression (3)

□ Thus the **gradient of the loss function** is:

$$\nabla L(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \phi_n$$



- ▶ due to the nonlinear logistic regression function it is not possible to find a closed-form solution
- ▶ however, the error function is **convex** and **gradient-based optimization** can be applied (also in an **online learning setting**)

# Multiclass Logistic Regression $k > 2$

- In multiclass problems,  $p(C_k|\phi)$  is modeled by a **softmax** transformation of the output of  $K$  linear functions (one for each class):

$$p(C_k|\phi) = y_k(\phi) = \frac{\exp(\mathbf{w}_k^T \phi)}{\sum_j \exp(\mathbf{w}_j^T \phi)}$$

- As for the two-class logistic regression and assuming 1-of- $K$  encoding for the target, we can compute **likelihood** as:

$$p(\mathbf{T}|\Phi, \mathbf{w}_1, \dots, \mathbf{w}_K) = \prod_{n=1}^N \underbrace{\left( \prod_{k=1}^K p(C_k|\phi_n)^{t_{nk}} \right)}_{\text{Only one term corresponding to correct class}} = \prod_{n=1}^N \left( \prod_{k=1}^K y_{nk}^{t_{nk}} \right)$$

$\downarrow$   
 $\mathbf{T} = (t_1, t_2, \dots, t_N)$

# Multiclass Logistic Regression

- In multiclass problems,  $p(C_k|\phi)$  is modeled by a **softmax** transformation of the output of K linear functions (one for each class):

$$p(C_k|\phi) = y_k(\phi) = \frac{\exp(\mathbf{w}_k^T \phi)}{\sum_j \exp(\mathbf{w}_j^T \phi)}$$

- As for the two-class logistic target, we can compute likelihood as.

$$y_{nk} = p(C_k|\phi_n) = \frac{\exp(\mathbf{w}_k^T \phi_n)}{\sum_j \exp(\mathbf{w}_j^T \phi_n)}$$

$$p(\mathbf{T}|\Phi, \mathbf{w}_1, \dots, \mathbf{w}_K) = \prod_{n=1}^N \underbrace{\left( \prod_{k=1}^K p(C_k|\phi_n)^{t_{nk}} \right)}_{\text{Only one term corresponding to correct class}} = \prod_{n=1}^N \left( \prod_{k=1}^K y_{nk}^{t_{nk}} \right)$$

Only one term corresponding  
to correct class

## Multiclass Logistic Regression (2)


- As for the two-class problem, we can minimize the cross-entropy error function:

$$L(\mathbf{w}_1, \dots, \mathbf{w}_K) = -\ln p(\mathbf{T}|\Phi, \mathbf{w}_1, \dots, \mathbf{w}_K) = -\sum_{n=1}^N \left( \sum_{k=1}^K t_{nk} \ln y_{nk} \right)$$


- Then, we compute the gradient for each weights vector:

$$\nabla L_{\mathbf{w}_j}(\mathbf{w}_1, \dots, \mathbf{w}_K) = \sum_{n=1}^N (y_{nj} - t_{nj}) \phi_n$$

- If we replace the logistic function with a step function...



$$y(\mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \phi(\mathbf{x})}}$$



$$y(\mathbf{x}, \mathbf{w}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \phi(\mathbf{x}) > 0 \\ 0 & \text{otherwise} \end{cases}$$

- ... logistic regression leads to the same **updating rule** of the perceptron algorithm:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha (y(\mathbf{x}_n, \mathbf{w}) - t_n) \phi_n$$